



The Architecture of Consensus

Designing Reliable Distributed Systems in an Unreliable World



The Impossibility of Perfect Certainty




No Common Knowledge

A node can never be mathematically certain of a remote node's exact state.

The Communication Paradox

Knowledge requires communication, but communication links are inherently unreliable. Certainty requires infinite message loops.

Designing for Malicious Actors



Rule of Tolerance:
$$N \geq 3f + 1$$

The Trust Challenge

Nodes may not merely crash—they may actively lie, collude, or deviate from protocols.

The Mathematical Limit

To tolerate f malicious nodes, a system strictly requires at least $3f + 1$ total nodes (fewer than 1/3 can be malicious).

The Baseline Assumptions of Distributed Environments

Network Behavior

Reliable

Perfect delivery (No loss).

Fair-Loss

Messages may be lost/delayed, but eventually arrive if retried.

Arbitrary

Active adversary can intercept, modify, or drop messages.

Node Behavior

Crash-Stop

Node fails and never returns.

Crash-Recovery

Node fails, loses memory, but resumes later using disk data.

Byzantine

Node acts maliciously or arbitrarily.

Takeaway: Every distributed algorithm is engineered against a specific combination of these exact constraints.

Securing Local Concurrency Before Network Coordination



Safety

Prevent buffer overflow (Producer blocks) and underflow (Consumer blocks).

Mechanism

Handled via counting semaphores or condition variables to ensure strict mutual exclusion.

The Foundational Principles of Global State



Atomicity

All or nothing. No partial or intermediate updates across the network.



Consistency

Valid states only. Global invariants are strictly maintained.



Isolation


Sequential illusion. Concurrent operations never interfere.



Durability

Permanent survival. Committed data survives crashes via write-ahead logs.

Wrapping the Network in a Single Operation




Spans multiple geographically separated hardware systems.

Requires a dedicated Coordinator to orchestrate outcomes.

If any single component fails, the entire distributed state must globally rollback.

Forcing Agreement with Two-Phase Commit



The Mechanism

A synchronous vote followed by a global execution command.


The Vulnerability

If the Coordinator crashes after participants vote Yes, participants are trapped waiting. Crucial resources remain locked indefinitely.

Escaping the Block with Three-Phase Commit


Phase 1 (Prepare)

Vote Yes/No




Phase 2 (Pre-Commit)

Intermediate state. Coordinator broadcasts intent.



Phase 3 (Commit)

Final execution.



The Fix

By broadcasting a Pre-Commit intent, nodes possess shared knowledge of the impending outcome.


The Result

If the Coordinator fails, participants can safely time out and reach consensus using shared logs—preventing infinite resource locks.

Selecting the Right Commit Strategy

	Two-Phase Commit (2PC)	Three-Phase Commit (3PC)
Failure Risk	Susceptible to indefinite blocking	Non-blocking (handles coordinator crash)
Network Overhead	Lower (2 message rounds)	Higher (3 message rounds)
Partitions	Vulnerable	Handles partitions gracefully
Ideal Use Case	Simple, highly reliable internal networks	Resilient, distributed workflows

Managing Network Uncertainty



Timeouts

Allow nodes to actively detect an unresponsive coordinator.

Exponential Backoff

Prevents network congestion by deliberately spacing out retry attempts.

Hard Limits

Essential cut-offs to prevent infinite loops and forcefully release locked resources.

Adapting Strict ACID for Modern Microservices



1	2	3
Strict global ACID (like 2PC) creates massive bottlenecks in independent microservice databases.	The Saga Pattern : Services execute sequentially and communicate asynchronously.	If a downstream service fails, the system executes Compensating Transactions to logically undo previous steps rather than holding hard global locks.

The Ultimate Metric: High Availability

The Nines of Availability

Availability	Concept	Downtime Allowed per Year
99%	Two Nines	3.7 Days
99.9%	Three Nines	8.8 Hours
99.99%	Four Nines	53 Minutes
99.999%	Five Nines	5.3 Minutes

Fault-tolerant commit protocols, exponential backoffs, and sagas are the architectural prerequisites to achieving Four or Five Nines on inherently unreliable hardware.

The Blueprint for Distributed Trust




Node

+



Network

+



Algorithm

Perfect certainty is impossible:

The Two Generals prove we operate in an environment of permanent doubt.

Protocols create order:

Algorithms like 2PC and 3PC force atomic outcomes across geographic divides.



Consistency
Shield

Models dictate design:

We must engineer specifically for fair-loss networks and crash-recovery nodes.

Resilience requires strategy:

Distributed safety relies entirely on engineered fallbacks: timeouts, retries, and compensation.